

Text Classification with Scikit-Learn

Objective

Build and evaluate a machine learning model that classifies text documents into categories. You will apply text preprocessing, feature extraction, model training, and performance evaluation using Python and scikit-learn. Dataset Options

Students use the following datasets:

1. 20 Newsgroups dataset and classify news posts into topics like politics, science, or sports.

```
from sklearn.datasets import fetch_20newsgroups
```

Tasks

- 1. Data Loading
 - Load the dataset
 - Print sample texts and their labels.
- 2. Text Preprocessing
 - Convert text to lowercase.
 - Remove punctuation and stopwords.
 - Apply stemming and lemmatization.
- 3. Feature Extraction
 - Use TfidfVectorizer and CountVectorizer from scikit-learn.
 - Optional: how n-grams affect model performance (e.g., unigram vs bigram).
`vectorizer = TfidfVectorizer(max_features=5000, ngram_range=(1,2))`
`X = vectorizer.fit_transform(texts)`
`ngram_range=(1,2) → Unigrams + bigrams` (word pairs, e.g., “not good”).
- 4. Model Training
 - Train at least two classifiers (e.g., MultinomialNB, LinearSVC, LogisticRegression).
 - Split dataset into training/test sets.
- 5. Model Evaluation
 - Compute and display: Accuracy, Precision, Recall, F1-score.
 - Plot confusion matrix using matplotlib.
 - Compare model results and discuss findings.

- 6. Experimentation
 - Optional: Try dimensionality reduction (e.g., TruncatedSVD for LSA).
 $svd = TruncatedSVD(n_components=100,...)$
 - Use a Pipeline to combine preprocessing + model steps.
 - Perform hyperparameter tuning with GridSearchCV.

Expected Output

- Printed metrics and plots.
- Explanation of preprocessing decisions.
- Discussion on model performance and insights from misclassified examples.

Submission Requirements

- Submit a single Jupyter Notebook (.ipynb).

Lab_3-TextClass-FirstName.ipynb

- Include Markdown explanations and code comments.
- Provide plots and a short summary paragraph at the end.